

Optimization in Neural Networks

DRIPTA MJ

Department of Mathematics

RAMAKRISHNA MISSION VIVEKANANDA EDUCATIONAL AND RESEARCH INSTITUTE
BELUR MATH, INDIA

Machine Learning
CS230

Sem 3, 2018-19

Gradient based optimization

- Gradient descent:

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \xi \sum_{n=1}^N \nabla_{\mathbf{w}^{(t)}} L(y^{(n)}, y^{*(n)}(\mathbf{w}^{(t)}))$$

where ξ is the learning rate.

- Frequency of updates:
 - Batch gradient descent: Updates after evaluating the loss gradient w.r.t. all training examples.
 - Stochastic gradient descent: Updates after evaluating the loss gradient w.r.t. every training example.
 - Mini-batch gradient descent: Updates after evaluating the loss gradient w.r.t. a subset of the training dataset.
- Type of updates:
 - Fixed learning rate
 - With momentum
 - Adaptive learning rate
 - Adaptive learning rate + Momentum

Ravines

- Stochastic gradient descent has difficulty navigating ravines.



Momentum based gradient descent

- Builds up speed in directions with gentle and consistent gradient.
- Damps oscillations in direction of high curvature.
- The effect of the gradient is to increment the previous velocity. The velocity also decay by a factor β which is slightly less than one.
- Running average makes the gradient less dependent on its current value, and rely more on the general behaviour of the gradient in the past updates.
- More interested in the expected value of the gradient rather on the particular gradient value at a particular iteration.

$$\mathbf{v}^{(t)} = \beta \mathbf{v}^{(t-1)} + (1 - \beta) \nabla_{\mathbf{w}^{(t)}} L$$

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \xi \mathbf{v}^{(t)}$$

Momentum based gradient descent: updates

$$\mathbf{v}^{(0)} = 0$$

- Update 1: $\mathbf{v}^{(1)} = \beta \mathbf{v}^{(0)} + (1 - \beta) \nabla_{\mathbf{w}^{(1)}} L$
 $= (1 - \beta) \nabla_{\mathbf{w}^{(1)}} L$
- Update 2: $\mathbf{v}^{(2)} = \beta \mathbf{v}^{(1)} + (1 - \beta) \nabla_{\mathbf{w}^{(2)}} L$
 $= \beta(1 - \beta) \nabla_{\mathbf{w}^{(1)}} L + (1 - \beta) \nabla_{\mathbf{w}^{(2)}} L$
- Update 3: $\mathbf{v}^{(3)} = \beta \mathbf{v}^{(2)} + (1 - \beta) \nabla_{\mathbf{w}^{(3)}} L$
 $= \beta(\beta(1 - \beta) \nabla_{\mathbf{w}^{(1)}} L + (1 - \beta) \nabla_{\mathbf{w}^{(2)}} L) + (1 - \beta) \nabla_{\mathbf{w}^{(3)}} L$
 $= (1 - \beta) [\beta^2 \nabla_{\mathbf{w}^{(1)}} L + \beta^1 \nabla_{\mathbf{w}^{(2)}} L + \nabla_{\mathbf{w}^{(3)}} L]$
- Update t : $\mathbf{v}^{(t)} = (1 - \beta) [\beta^{(t-1)} \nabla_{\mathbf{w}^{(1)}} L + \beta^{(t-2)} \nabla_{\mathbf{w}^{(2)}} L + \dots + \nabla_{\mathbf{w}^{(t)}} L]$

Shortcomings

- Binary classification problem
- Binary cross-entropy loss function

- Update: $w_i := w_i - \xi \frac{\partial L}{\partial w_i}$


- The gradient can be computed as $\frac{\partial L}{\partial w_i} = (y^* - y)x_i$

- Suppose the i -th feature is sparse, i.e. $x_i = 0$ for most of the examples.

- In that case the weight w_i associated with x_i will have few updates as the gradient is 0 in most cases.

- Our results can be seriously impacted if x_i happens to be a very important feature.

- Want an algorithm that gives higher learning rate to sparse features.



Adagrad

- Adapts the learning rate of the parameters.
 - Higher learning rate for sparse features.
 - Lower learning rate for dense features.
- More updates of a parameter indicate more decay of its learning rate.

$$\begin{aligned} \mathbf{s}^{(t)} &= \mathbf{s}^{(t-1)} + (\nabla_{\mathbf{w}^{(t)}} L)^2 \\ \mathbf{w}^{(t+1)} &= \mathbf{w}^{(t)} - \frac{\xi}{\sqrt{\mathbf{s}^{(t)} + \epsilon}} \odot \nabla_{\mathbf{w}^{(t)}} L \end{aligned}$$

- Method appropriate for dealing with sparse datasets.

RMSProp

- Adagrad reduces the learning rates of parameters associated with dense features very fast. So the corresponding weight updates will be small.
- **Adagrad**: Sum of squares of the past gradients.
- **RMSPprop**: Exponentially decaying (moving) average of the squares of past gradients.

$$\begin{aligned} \mathbf{s}^{(t)} &= \beta \mathbf{s}^{(t-1)} + (1 - \beta) (\nabla_{\mathbf{w}^{(t)}} L)^2 \\ \mathbf{w}^{(t+1)} &= \mathbf{w}^{(t)} - \frac{\xi}{\sqrt{\mathbf{s}^{(t)} + \epsilon}} \odot \nabla_{\mathbf{w}^{(t)}} L \end{aligned}$$

- RMSPprop addresses the learning rate decay problem of Adagrad.

ADAM

- Adaptive Moment Estimation (ADAM).
- Adaptive learning rate + Momentum
 - Keeps an exponentially decaying average of past gradients (like momentum).
 - Also keeps an exponentially decaying average of past squared gradients (like RMSProp).

$$\begin{aligned}\mathbf{v}^{(t)} &= \beta_1 \mathbf{v}^{(t-1)} + (1 - \beta_1) \nabla_{\mathbf{w}^{(t)}} L \\ \mathbf{s}^{(t)} &= \beta_2 \mathbf{s}^{(t-1)} + (1 - \beta_2) (\nabla_{\mathbf{w}^{(t)}} L)^2\end{aligned}$$

- $\mathbf{v}^{(t)}$ is the vector of first moment (mean) estimates of the mean of the gradients.
- $\mathbf{s}^{(t)}$ is the vector of second moment (uncentered variance) estimates of the gradients.
- Bias corrections: $\bar{\mathbf{v}}^{(t)} = \frac{\mathbf{v}^{(t)}}{1 - \beta_1^t}$ $\bar{\mathbf{s}}^{(t)} = \frac{\mathbf{s}^{(t)}}{1 - \beta_2^t}$
- Update of weights: $\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \frac{\xi}{\sqrt{\bar{\mathbf{s}}^{(t)} + \epsilon}} \odot \bar{\mathbf{v}}^{(t)}$

ADAM: Bias correction

- We have seen that $\mathbf{v}^{(t)} = (1 - \beta_1) \sum_{j=1}^t \beta^{t-j} \nabla_{\mathbf{w}^{(j)}} L$

- Taking Expectation on both sides yields

$$\begin{aligned}\mathbb{E}[\mathbf{v}^{(t)}] &= \mathbb{E}\left[(1 - \beta_1) \sum_{j=1}^t \beta^{t-j} \nabla_{\mathbf{w}^{(j)}} L\right] \\ &= (1 - \beta_1) \mathbb{E}\left[\sum_{j=1}^t \beta^{t-j} \nabla_{\mathbf{w}^{(j)}} L\right] \\ &= (1 - \beta_1) \sum_{j=1}^t \mathbb{E}\left[\beta^{t-j} \nabla_{\mathbf{w}^{(j)}} L\right] \\ &= (1 - \beta_1) \sum_{j=1}^t \beta^{t-j} \mathbb{E}\left[\nabla_{\mathbf{w}^{(j)}} L\right]\end{aligned}$$

- Assumption: All gradients come from the same distribution, i.e.

$$\nabla_{\mathbf{w}^{(1)}} L = \nabla_{\mathbf{w}^{(2)}} L = \dots\dots\dots = \nabla_{\mathbf{w}^{(j)}} L = \nabla_{\mathbf{w}} L$$

ADAM: Bias correction

- So then we have

$$\begin{aligned}\mathbb{E}[\mathbf{v}^{(t)}] &= (1 - \beta_1) \sum_{j=1}^t \beta_1^{t-j} \mathbb{E}[\nabla_{\mathbf{w}} L] \\ &= (1 - \beta_1) \mathbb{E}[\nabla_{\mathbf{w}} L] \sum_{j=1}^t \beta_1^{t-j} \\ &= (1 - \beta_1) \mathbb{E}[\nabla_{\mathbf{w}} L] (\beta_1^{t-1} + \beta_1^{t-2} + \dots + \beta_1^0) \\ &= (1 - \beta_1) \mathbb{E}[\nabla_{\mathbf{w}} L] \left(\frac{1 - \beta_1^t}{1 - \beta_1} \right) \\ &= \mathbb{E}[\nabla_{\mathbf{w}} L] (1 - \beta_1^t)\end{aligned}$$

Therefore

$$\begin{aligned}\mathbb{E}\left[\frac{\mathbf{v}^{(t)}}{1 - \beta_1^t}\right] &= \mathbb{E}[\nabla_{\mathbf{w}} L] \\ \mathbb{E}[\bar{\mathbf{v}}^{(t)}] &= \mathbb{E}[\nabla_{\mathbf{w}} L]\end{aligned}$$

Comparison




Figure source: Kingma and Lei Ba, ADAM: A method for stochastic optimization, 2015.