

Regularization

DRIPTA MJ


Department of Mathematics

RAMAKRISHNA MISSION VIVEKANANDA EDUCATIONAL AND RESEARCH INSTITUTE
BELUR MATH, INDIA

Machine Learning
CS230

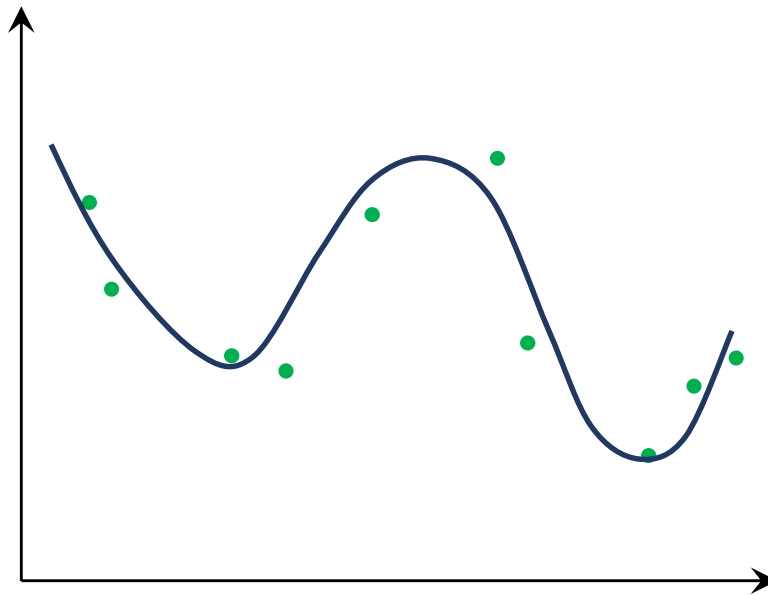
Sem 3, 2018-19

Bias-Variance trade-off



Polynomial curve fitting

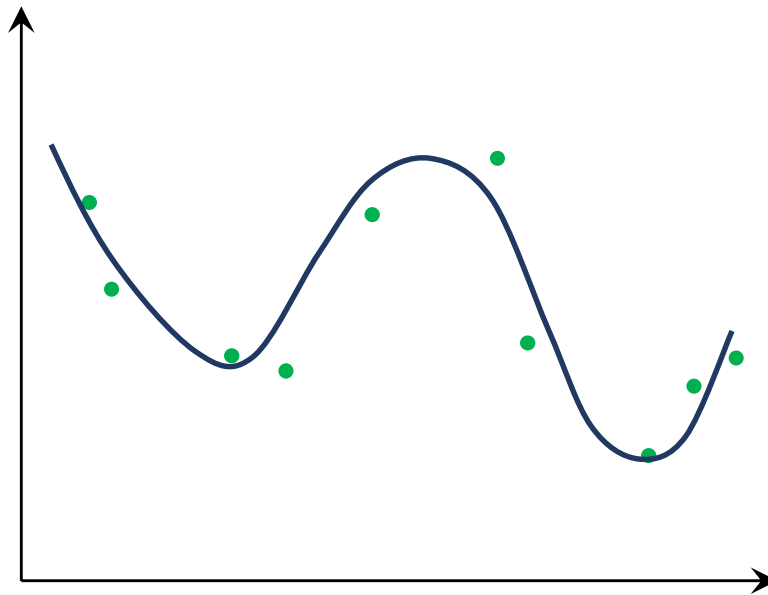
- Data generated by a Q th order polynomial + some noise



- Consider fitting data with a polynomial of order M .
- Data preprocessing:
 - Standardize the inputs.
 - Center the outputs.
- Model can be trained using linear regression with $[x^1, x^2, \dots, x^M]$ as features.
- The intercept w_0 can then be ignored.

Polynomial curve fitting

- Data generated by a Q th order polynomial + some noise

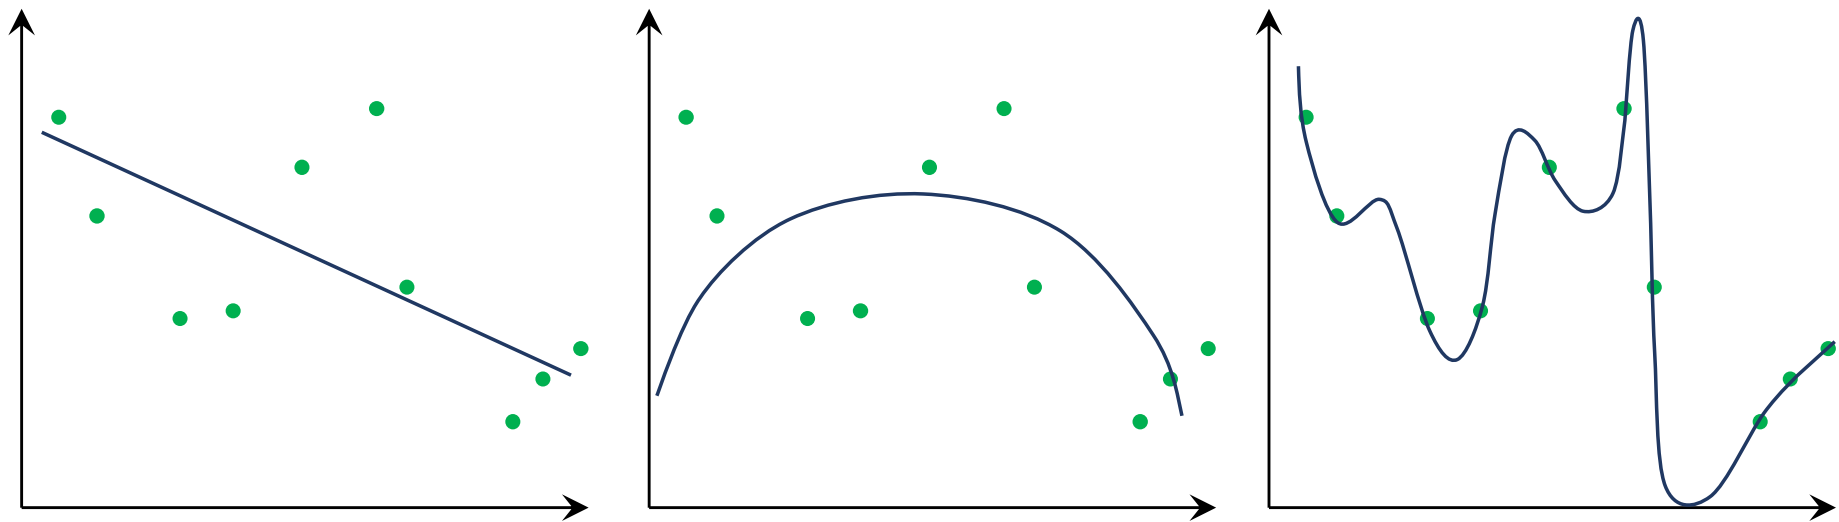


- Predictor model:

$$\begin{aligned} f(x, \mathbf{w}) &= w_1x + w_2x^2 + w_3x^3 + \dots + w_Mx^M \\ &= \sum_{i=1}^M w_i x^i \\ &= \mathbf{w}^T \phi \end{aligned}$$

where $\mathbf{w} = [w_1, \dots, w_M]^T$ and $\phi = [x, \dots, x^M]^T$.

Polynomial curve fitting



- Complex hypotheses (richer class of models) lead to overfitting.
- A higher degree polynomial has more degrees of freedom which can lead to overfitting of the training data.
- Need to penalize the complexity in some way in the cost function.

*Figures are just for illustration.

Regularized regression

- Observations:
 - Weights \mathbf{w} are unconstrained, and as such can lead to high variance.
 - Need to control the magnitude of the weights in order to control the variance.

- Modified objective:

$$\text{minimize} \quad \sum_{i=1}^N \left(y^{(i)} - \mathbf{w}^T \phi(\mathbf{x}^{(i)}) \right)^2 \quad \text{such that} \quad \sum_{j=1}^M w_j^2 \leq p$$

- In vector form:

$$\text{minimize} \quad (\mathbf{y} - \Phi \mathbf{w})^T (\mathbf{y} - \Phi \mathbf{w}) \quad \text{such that} \quad \|\mathbf{w}\|_2^2 \leq p$$

- Assumptions:

Φ is standardized (zero mean and unit variance), and \mathbf{y} is centered.

Regularized regression

- Can show that the problem is equivalent to:

$$\text{minimize} \quad (\mathbf{y} - \Phi \mathbf{w})^T (\mathbf{y} - \Phi \mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

where λ is the regularization coefficient.

- λ tries to balance between the fit to the training data and the model complexity.
- Modified loss function:

$$L(\mathbf{w}) = L_E(\mathbf{w}) + \lambda L_R(\mathbf{w})$$

where

$$\begin{aligned} L_E(\mathbf{w}) &= \sum_{i=1}^N (y^{(i)} - \mathbf{w}^T \phi(\mathbf{x}^{(i)}))^2 \\ &= (\mathbf{y} - \Phi \mathbf{w})^T (\mathbf{y} - \Phi \mathbf{w}) \end{aligned}$$

$$\begin{aligned} L_R(\mathbf{w}) &= \sum_{j=1}^M w_j^2 \\ &= \|\mathbf{w}\|_2^2 \end{aligned}$$

L_2 regularization


- This is known as **L_2 Regularization** and also as **Ridge Regression**.
- Goal is to minimize the loss function $L(\mathbf{w})$. Note, since $L(\mathbf{w})$ is convex it has a unique solution.
- Taking derivative of $L(\mathbf{w})$ with respect to \mathbf{w} and equating it to zero

$$\left(\text{i.e. } \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = 0 \right) \text{ yields:}$$

$$\mathbf{w} = (\Phi^T \Phi + \lambda \mathbf{I})^{-1} \Phi^T \mathbf{y}$$

- If $\lambda = 0$, we get the least squares solutions.
- If $\lambda \rightarrow \infty$, we get $\mathbf{w} \rightarrow 0$.
- So $\lambda > 0$ will give weights of lower magnitudes than that obtained using least squares.

Visualization (L_2)



L_1 regularization

- Use L_1 norm of the weight vector.

$$\text{minimize} \quad (\mathbf{y} - \Phi \mathbf{w})^T (\mathbf{y} - \Phi \mathbf{w}) \quad \text{such that} \quad \sum_{j=1}^M |w_j| \leq p$$

- Known as the **LASSO** (least absolute shrinkage and selection operator) algorithm (*Tibshirani*, 1996).
- **LASSO** has no closed form solution unlike ridge regression.
- Can be solved using quadratic programming techniques.
- Often want some of the weights w_j 's to be 0.
- **LASSO** looks for a sparse solution and so likely to yield some of the weights to be 0. But why?

L_1 regularization

- Consider a problem with two features x_1 and x_2 .
- In this case we are trying to solve a optimization problem with respect to weights w_1 and w_2 :

$$\text{minimize} \quad \sum_{i=1}^N \left(y^{(i)} - w_1 x_1^{(i)} - w_2 x_2^{(i)} \right)^2$$

such that


$$w_1 + w_2 \leq p$$

$$-w_1 + w_2 \leq p$$

$$w_1 - w_2 \leq p$$

$$-w_1 - w_2 \leq p$$

Visualization (L_1)



- When λ is large, then among the contours satisfying the constraints, the contour with the least value of the objective function is likely to intersect the constraints' boundary at a corner.

K-fold cross validation

- Training data is subdivided into K separate subsets – $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_K$ of equal size (say n_K).
- For $k = 1, 2, \dots, K$
 - Leave out the k th fold data \mathcal{D}_k and train the model on the remaining $k - 1$ folds.
 - Use the trained model to make prediction on the k th fold data \mathcal{D}_k and compute the (cross validation) error for this fold

$$E_k^{(\lambda)} = \frac{1}{n_K} \sum_{i=1}^{n_K} (y_{k,i} - f_{-k}^{(\lambda)}(\mathbf{x}_i))^2$$

where $f_{-k}^{(\lambda)}$ is the model trained excluding the k th fold data with a specific value of λ .

K-fold cross validation

- Estimated generalization error:

$$\mathbf{E}^{(\lambda)} = \frac{1}{K} \sum_{k=1}^K E_k^{(\lambda)}$$

- The optimal value of λ (say λ^*) is the one yielding the least value of $\mathbf{E}^{(\lambda)}$.

- Using λ^* train the model on the entire training dataset.
- When $K = N$ (size of the training dataset), the approach is known as [leave-one-out cross-validation](#).